


Class begins on Michigan Time. Before class begins, please do the following:


1. Visit [www.umich.edu/~cja/](http://www.umich.edu/~cja/) and click "STATS 531 Introduction to Flux" to view these slides.
2. Login to Flux using the instructions on the first few slides.



# STATS 531

## Introduction to Flux

Dr Charles J Antonelli, LSA IT ARS  
Mark Montague, LSA IT ARS  
February, 2018



# STATS 531

## Introduction to Flux

Dr Charles J Antonelli, LSA IT ARS  
Mark Montague, LSA IT ARS  
March, 2018

# Preliminaries

# Connecting to Flux

## (Windows)

1. If you don't have PuTTY installed  
Install the PuTTY/WinSCP installer for Windows:  
Compute at the U: <http://its.umich.edu/computing/computers-software/compute>  
Select + under Get Going  
Get U-M PuTTY (UM\_PuTTY\_WinSCP.zip)  
Execute the installer (UM\_PuTTY\_WinSCP.exe)  
Accept all defaults
2. If you don't see this icon in your system tray  
Install the Xming X Server for Windows  
<https://sourceforge.net/projects/xming/files/Xming/6.9.0.31/>  
  
Get the Xming installer (Xming-6-9-0-31-setup.exe)  
Execute the installer  
Accept all defaults  
<https://sourceforge.net/projects/xming/files/Xming-fonts/7.7.0.10/>  
Get the Xming fonts installer (Xming-fonts-7-7-0-10-setup.exe)  
Execute the installer  
Accept all defaults

# Connecting to Flux

## (Windows)

1. Double-click "UM Internet Access Kit" icon on Desktop  
Double-click PuTTY application within
2. In the Putty Application that appears:  
Enter "flux-login.arc-ts.umich.edu" in the Host Name box  
Under Connection | SSH | X11, ensure Enable X11 Forwarding is checked  
Click "Open" at bottom
3. In the terminal window that appears:  
Login with username, Kerberos password, and Duo
4. This creates an ssh session on Flux

# Connecting to Flux

(Linux & macOS)

1. Start a terminal window
2. In the terminal window that appears, type  
`ssh -X username@flux-login.arc-ts.umich.edu`
3. Login to Flux with Kerberos password and Duo



# Roadmap

- Preliminaries
- Introduction to Flux
- The command line
- Flux mechanics
- Parallel R and [pomp](#)

# Introduction to Flux



# Flux

Flux is a university-wide shared computational discovery / high-performance computing service.


- Provided by Advanced Research Computing at U-M
- Procurement, licensing, billing by U-M ITS
- Interdisciplinary since 2010

<http://arc-ts.umich.edu/resources/compute-resources/>



# The Flux cluster

Login nodes




Data transfer  
node



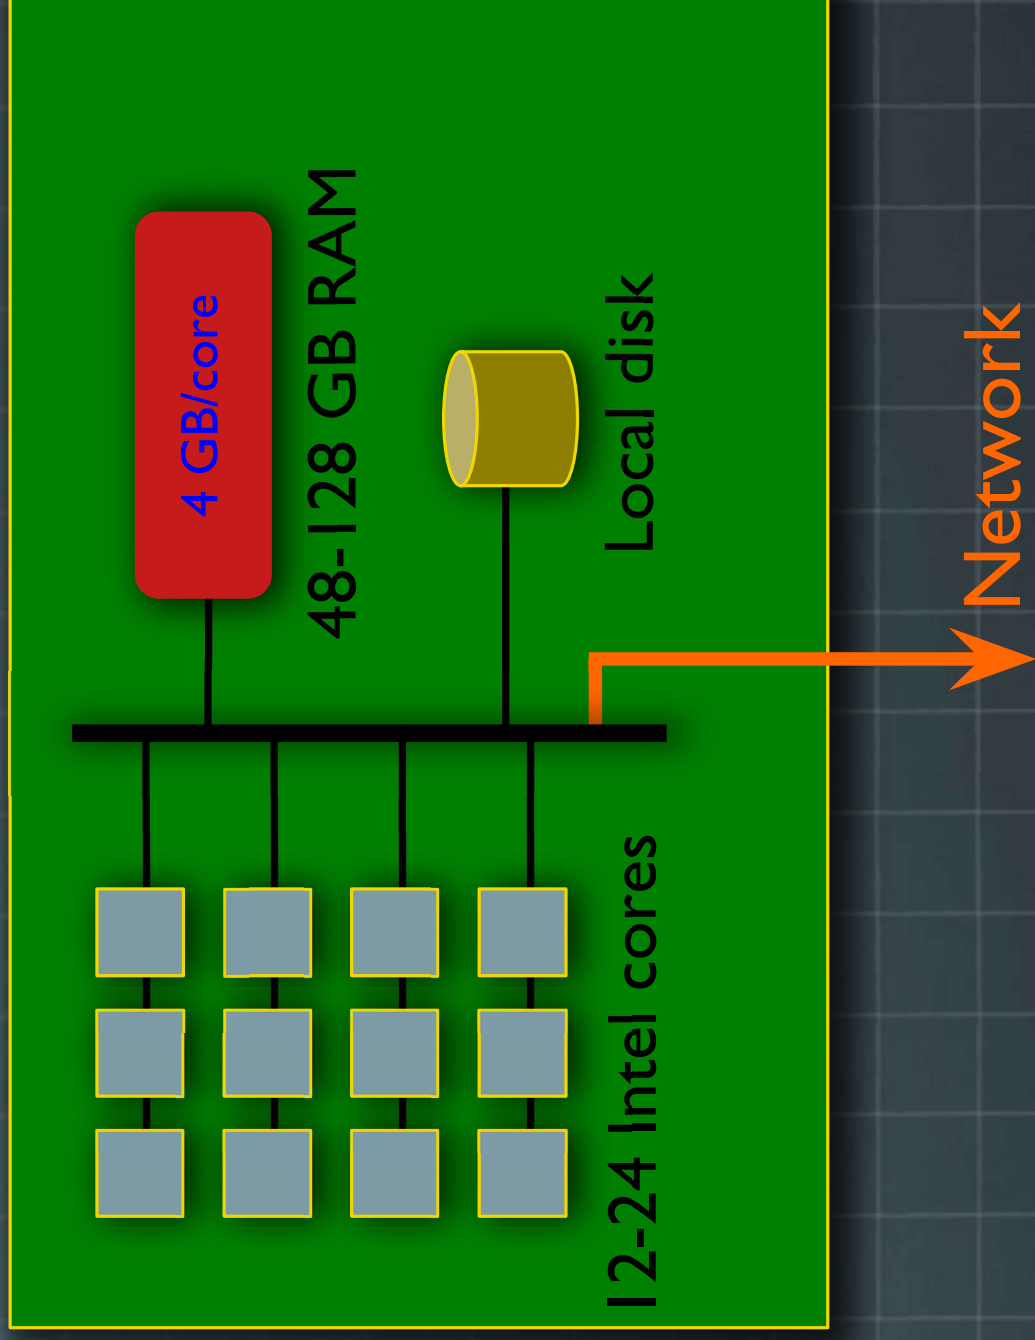
Storage



Compute nodes



# A Standard Flux node



# Using Flux

- Three basic requirements:
  - A Flux login account  
<https://arc-ts.umich.edu/fluxform>
  - A Flux allocation  
[stats531w18\\_flux](#)
  - A Duo app on your smartphone or tablet  
<http://its.umich.edu/two-factor-authentication>
- Logging in to Flux
  - `ssh -X login@flux-login.arc-ts.umich.edu`
  - PuTTY
  - Campus wired or MWireless
  - Otherwise use VPN, or
  - `ssh login.itd.umich.edu` first

# The command line




# Command Line Reference

William E Shotts, Jr.,  
“The Linux Command Line: A Complete Introduction,”  
No Starch Press, January 2012.  
<http://linuxcommand.org/tlcl.php> .

Download Creative Commons Licensed  
version at

[http://downloads.sourceforge.net/project/  
linuxcommand/TLCL/13.07/TLCL-13.07.pdf](http://downloads.sourceforge.net/project/linuxcommand/TLCL/13.07/TLCL-13.07.pdf)





# At the command prompt

- Basic input line editing commands
  - **Backspace** erases previous character
  - **Left** and **right arrow** move insertion point on the line
  - **Control-U** erases the line to the insertion point, so you can start over
  - **Enter** executes the line you typed
  - **Control-C** interrupts whatever command you started and returns you to the shell prompt (usually)
  - **Up** and **down arrow** will access your *command history*
  - Type “exit” without the quotes to exit the shell

# Flux mechanics

# Cluster batch workflow

- You create a batch script and submit it to PBS
- PBS schedules your job, and it enters the flux queue
- When its turn arrives, your job will execute the batch script
- Your script has access to all Flux applications and data
- When your script completes, anything it sent to standard output and error are saved in files stored in your submission directory
- You can ask that email be sent to you when your jobs starts, ends, or aborts
- You can check on the status of your job at any time, or delete it if it's not doing what you want
- A short time after your job completes, it disappears from PBS

# Multi-threaded batch script

```
#PBS -N yourjobname
#PBS -V
#PBS -A youralloc_flux
#PBS -q flux
#PBS -l nodes=1:ppn=12,mem=47gb,walltime=00:05:00
#PBS -m abe
#PBS -j oe
```

```
#Your Code Goes Below:
cd $PBS_O_WORKDIR
R CMD BATCH --vanilla myscript.R myscript.out
```

# Copying data

## Using command line programs:

scp: copies files between hosts on a network over ssh  
`scp localfile username@flux-xfer.arc-ts.umich.edu:remotefile`  
`scp -r localdir username@flux-xfer.arc-ts.umich.edu:remotedir`  
`scp username@flux-login.arc-ts.umich.edu:remotefile localfile`

Use "." as destination to copy to your Flux home directory:

`scp localfile username@flux-xfer.arc-ts.umich.edu:.`

... or to your Flux scratch directory:

`scp localfile username@flux-xfer.arc-ts.umich.edu:/scratch/allcname/username`

sftp: an interactive file transfer program over ssh (a secure ftp)  
`sftp username@flux-xfer.arc-ts.umich.edu`

## Using graphical (GUI) applications:

FileZilla (cross-platform):

<http://filezilla-project.org/>

Cyberduck (Mac):

<https://cyberduck.io/>

WinSCP (Windows):

<http://www.itcs.umich.edu/bluedisc/>



# Basic batch commands

- Once you have a script, submit it:  
`qsub scriptfile`

```
$ qsub singlenode.pbs
6023521.nyx.engin.umich.edu
```

- You can check on the job status  
`qstat jobid`  
`qstat -u user`

```
$ qstat -u cja
nyx.engin.umich.edu:
```

Job ID	Username	Queue	Jobname	SessID	NDS	TSK	Memory	Req'd Time	Req'd S	Elap Time
6023521.nyx.engi	cja	flux	hpc101i	--	1	1	--	00:05	Q	--

- To delete your job  
`qdel jobid`



# Modules

- The module command allows you to specify what Flux software you want to use

```
module list
module load name
module avail
module avail name
module key string
module spider string
module unload name
module use path
module

-- Show loaded modules
-- Load module name for use
-- Show all available modules
-- Show versions of module name
-- Search for string in module descrip
-- Search for string in all module doc
-- Unload module name
-- Add path to module search path
-- List all options
```
- Enter these commands at any time during your session
- You can load multiple modules in the same login session
- Software modules remain available throughout the session

# Modules

- Once loaded, you can group a set of modules into a `module set`  
`module save myset`
- A module set can be restored at any time  
`module restore myset`
- List all module sets you've defined  
`module saveList`
- Create a module set to be loaded each time you log in  
`module save`

# Lab

Task: Use the R multicore package

- Copy sample code to your login directory

```
cd
```

```
cp
```

```
/scratch/data/workshops/stats/stats-sample-  
code.tar.gz .  
tar -zxvf stats-sample-code.tar.gz  
cd ./stats-sample-code
```

- Examine lab3.pbs and lab3.R

# Lab

Task: Use the R multicore package

- `module load R`
- Submit your job to Flux  
`qsub lab3.pbs`
- Watch the progress of your job  
`qstat -u username`  
where *username* is your own username
- When complete, look at the job's output  
`less lab3.out`

# Parallel R and pomp

# Random numbers

## Serial pseudorandom sequence

- Usually based on linear recurrences modulo  $m$
- Initialize the generator with some seed  $s$
- Generate stream of pseudorandom numbers

## Parallel pseudorandom sequence

- Performance, Reproducibility, Serializability



# Install pomp

Task: Install pomp package

- `module load R/3.4.1`
- `R`  
`install.packages("pomp", repos="https://`  
`/cran.mtu.edu/")`  
(answer y to "Would you like to use a personal library instead?" and "Would you like to create a personal library")

List of mirrors: <https://cran.r-project.org/mirrors.html>

# Multicore example

```
# To submit job: qsub pomp-par.pbs  
# Code shown below from: pomp-par.R
```

```
rm(list=ls())  
library(doParallel)  
set.seed(2018,kind="L'Ecuyer")  
cores <- as.numeric(Sys.getenv('PBS_NP', unset='8'))  
cl <- makeCluster(cores)  
registerDoParallel(cl)  
trials <- 100
```

# Multicore example

```
# pomp-par.R continued...
system.time(
  r <- foreach(icount(trials),
    .inorder=FALSE,
    .options.multicore=list(set.seed=TRUE)
  ) %dopar% {
  library(pomp)
  for (i in 1:100) {
    pompExample("gompertz")
    simulate(gompertz)
  }
}
)
stopCluster(cl)
save(r, file="sims.Rda")
```

# MPI example

- Change the R module:  
module unload R/3.4.1  
module load Rmpi/3.4.1
- Install doMPI package:  
Rmpi  
install.packages("doMPI", repos="https://cran.mtu.edu/")

# MPI example

```
# To submit job: qsub pomp-mpi.pbs
# Code shown below from: pomp-mpi.R
# changes from multicore example shown in green

rm(list=ls())
library(doMPI)
set.seed(2018, kind="L'Ecuyer")
cores <- as.numeric(Sys.getenv('PBS_NP', unset='8'))
cl <- startMPIcluster(cores)
registerDoMPI(cl)
trials <- 100
```

# MPI example

```
# pomp-mpi.R continued...
system.time(
  r <- foreach(icount(trials),
    .inorder=FALSE,
    .options.multicore=list(set.seed=TRUE)
  ) %dopar% {
  library(pomp)
  for (i in 1:100) {
    pompExample("gompertz")
    simulate(gompertz)
  }
}
)
closeCluster(cl)
save(r, file="sims.Rda")
```



# Interactive jobs

- You can submit jobs *interactively*:

```
qsub -I -X -V -l nodes=1:ppn=2 -l walltime=15:00  
-A youralloc_flux -l qos=flux -q flux
```

- This queues a job as usual
  - Your terminal session will be blocked until the job runs
  - When your job runs, you'll get an interactive shell
  - When you exit the shell your job is deleted
- Interactive jobs allow you to
  - Develop and test on cluster node(s)
  - Execute GUI tools on a cluster node
  - Utilize a parallel debugger interactively

# Lab

Task: Use an interactive PBS session

- `module load R`
- Start an interactive session  
`qsub -I -V -l nodes=1:ppn=2 -l walltime=30:00 -A stats531w18_flux -q flux`
- Run R in the interactive shell  
`cd $PBS_O_WORKDIR`  
`R`

# Troubleshooting

- System-level
  - `freenodes`
    - # aggregate node/core busy/free
    - # nodes, states, properties
    - # with `-l`, list only nodes marked down
  - `pbsnodes [-l]`
- Account-level
  - `mdiag -a acct`
    - # cores & users for account *acct*
  - `showq [-r][-i][-b][-w acct=acct]`
    - # running/idle/blocked jobs for *acct*
    - # with `-r|i|b` show more info for that job state
  - `freealloc [--jobs] acct`
    - # free resources in *acct*
    - # with `--jobs`, shows resources in use
  - `idlenodes acct [property]`
    - # shows available nodes for *acct* with property
- User-level
  - `mdiag -u uniq`
    - # allocations for user *uniq*
  - `showq [-r][-i][-b][-w user=uniq]`
    - # running/idle/blocked jobs for *uniq*
- Job-level
  - `qstat -f jobno`
    - # full info for job *jobno*
  - `qstat -n jobno`
    - # show nodes/cores where *jobno* running
  - `checkjob [-v] jobno`
    - # show why *jobno* not running
  - `qpeek jobno`
    - # peek at script output while *jobno* is running

# Resources

- ARC User Guide <http://arc-ts.umich.edu/flux-user-guide/>
- ARC Flux pages <http://arc-ts.umich.edu/flux/>
- Software Catalog <http://arc-ts.umich.edu/software/>
- Quick Start Guide <http://arc-ts.umich.edu/flux/using-flux/flux-in-10-easy-steps/>
- Flux FAQs <http://arc-ts.umich.edu/flux/flux-faqs/>
- For assistance, send email to: [hpc-support@umich.edu](mailto:hpc-support@umich.edu)
  - Read by a team of people including unit support staff
  - Can help with Flux operational and usage questions
  - Programming support available